



Red Hat OpenShift AI Self-Managed 3.0

Installing and uninstalling OpenShift AI Self-Managed

Install and uninstall OpenShift AI Self-Managed

Red Hat OpenShift AI Self-Managed 3.0 Installing and uninstalling OpenShift AI Self-Managed

Install and uninstall OpenShift AI Self-Managed

Legal Notice

Copyright © 2025 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

Install and uninstall OpenShift AI Self-Managed on your OpenShift cluster.

Table of Contents

PREFACE	4
CHAPTER 1. ARCHITECTURE OF OPENSIFT AI SELF-MANAGED	5
CHAPTER 2. UNDERSTANDING UPDATE CHANNELS	7
CHAPTER 3. INSTALLING AND DEPLOYING OPENSIFT AI	9
3.1. REQUIREMENTS FOR OPENSIFT AI SELF-MANAGED	9
3.1.1. Platform requirements	9
3.1.2. Component requirements	12
3.2. CONFIGURING CUSTOM NAMESPACES	13
3.3. INSTALLING THE RED HAT OPENSIFT AI OPERATOR	15
3.3.1. Installing the Red Hat OpenShift AI Operator by using the CLI	16
3.3.2. Installing the Red Hat OpenShift AI Operator by using the web console	18
3.4. INSTALLING AND MANAGING RED HAT OPENSIFT AI COMPONENTS	20
3.4.1. Installing Red Hat OpenShift AI components by using the CLI	20
3.4.2. Installing Red Hat OpenShift AI components by using the web console	23
3.4.3. Updating the installation status of Red Hat OpenShift AI components by using the web console	26
3.4.4. Viewing installed OpenShift AI components	29
CHAPTER 4. CONFIGURING PIPELINES WITH YOUR OWN ARGO WORKFLOWS INSTANCE	30
CHAPTER 5. INSTALLING THE DISTRIBUTED WORKLOADS COMPONENTS	32
CHAPTER 6. ACCESSING THE DASHBOARD	35
CHAPTER 7. ENABLING ACCELERATORS	36
CHAPTER 8. WORKING WITH CERTIFICATES	38
8.1. UNDERSTANDING HOW OPENSIFT AI HANDLES CERTIFICATES	38
8.2. ADDING CERTIFICATES	39
8.3. ADDING CERTIFICATES TO A CLUSTER-WIDE CA BUNDLE	40
8.4. ADDING CERTIFICATES TO A CUSTOM CA BUNDLE	41
8.5. USING SELF-SIGNED CERTIFICATES WITH OPENSIFT AI COMPONENTS	42
8.5.1. Accessing S3-compatible object storage with self-signed certificates	42
8.5.2. Configuring a certificate for pipelines	43
8.5.3. Configuring a certificate for workbenches	45
8.5.4. Using the cluster-wide CA bundle for the model serving platform	45
8.6. MANAGING CERTIFICATES WITHOUT THE RED HAT OPENSIFT AI OPERATOR	46
8.7. REMOVING THE CA BUNDLE	48
8.7.1. Removing the CA bundle from all namespaces	48
8.7.2. Removing the CA bundle from a single namespace	49
CHAPTER 9. VIEWING LOGS AND AUDIT RECORDS	50
9.1. CONFIGURING THE OPENSIFT AI OPERATOR LOGGER	50
9.1.1. Viewing the OpenShift AI Operator logs	51
9.2. VIEWING AUDIT RECORDS	51
CHAPTER 10. TROUBLESHOOTING COMMON INSTALLATION PROBLEMS	54
10.1. THE RED HAT OPENSIFT AI OPERATOR CANNOT BE RETRIEVED FROM THE IMAGE REGISTRY	54
10.2. OPENSIFT AI DOES NOT INSTALL ON UNSUPPORTED INFRASTRUCTURE	54
10.3. THE CREATION OF THE OPENSIFT AI CUSTOM RESOURCE (CR) FAILS	55
10.4. THE CREATION OF THE OPENSIFT AI NOTEBOOKS CUSTOM RESOURCE (CR) FAILS	55
10.5. THE OPENSIFT AI DASHBOARD IS NOT ACCESSIBLE	56

10.6. REINSTALLING OPENSIFT AI FAILS WITH AN ERROR	56
10.7. THE DEDICATED-ADMINS ROLE-BASED ACCESS CONTROL (RBAC) POLICY CANNOT BE CREATED	57
10.8. THE ODH PARAMETER SECRET DOES NOT GET CREATED	58
CHAPTER 11. UNINSTALLING RED HAT OPENSIFT AI SELF-MANAGED	59
11.1. UNDERSTANDING THE UNINSTALLATION PROCESS	59
11.2. UNINSTALLING OPENSIFT AI SELF-MANAGED BY USING THE CLI	60

PREFACE

Learn how to use both the OpenShift CLI (**oc**) and web console to install Red Hat OpenShift AI Self-Managed on your OpenShift cluster. To uninstall the product, learn how to use the recommended command-line interface (CLI) method.



IMPORTANT

You cannot upgrade from OpenShift AI 2.25 or any earlier version to 3.0. OpenShift AI 3.0 introduces significant technology and component changes and is intended for new installations only. To use OpenShift AI 3.0, install the Red Hat OpenShift AI Operator on a cluster running OpenShift Container Platform 4.19 or later and select the **fast-3.x** channel.

Support for upgrades will be available in a later release, including upgrades from OpenShift AI 2.25 to a stable 3.x version.

For more information, see the [Why upgrades to OpenShift AI 3.0 are not supported](#) Knowledgebase article.



NOTE

Red Hat does not support installing more than one instance of OpenShift AI on your cluster.

Red Hat does not support installing the Red Hat OpenShift AI Operator on the same cluster as the Red Hat OpenShift AI Add-on.

CHAPTER 1. ARCHITECTURE OF OPENSIFT AI SELF-MANAGED

Red Hat OpenShift AI Self-Managed is an Operator that is available in a self-managed environment, such as Red Hat OpenShift Container Platform, or in Red Hat-managed cloud environments such as Red Hat OpenShift Dedicated (with a Customer Cloud Subscription for AWS or GCP), Red Hat OpenShift Service on Amazon Web Services (ROSA classic or ROSA HCP), or Microsoft Azure Red Hat OpenShift.

OpenShift AI integrates the following components and services:

- At the service layer:

OpenShift AI dashboard

A customer-facing dashboard that shows available and installed applications for the OpenShift AI environment as well as learning resources such as tutorials, quick starts, and documentation. Administrative users can access functionality to manage users, clusters, workbench images, and model-serving runtimes. Data scientists can use the dashboard to create projects to organize their data science work.

Model serving

Data scientists can deploy trained machine-learning models to serve intelligent applications in production. After deployment, applications can send requests to the model using its deployed API endpoint.

AI pipelines

Data scientists can build portable machine learning (ML) workflows with AI pipelines by using Docker containers. With AI pipelines, data scientists can automate workflows as they develop their data science models.

Jupyter (self-managed)

A self-managed application that allows data scientists to configure a basic standalone workbench and develop machine learning models in JupyterLab.

Distributed workloads

Data scientists can use multiple nodes in parallel to train machine-learning models or process data more quickly. This approach significantly reduces the task completion time, and enables the use of larger datasets and more complex models.

Retrieval-Augmented Generation (RAG)

Data scientists and AI engineers can leverage Retrieval-Augmented Generation (RAG) capabilities provided by the integrated Llama Stack Operator. By combining large language model inference, semantic retrieval, and vector database storage, data scientists and AI engineers can obtain tailored, accurate, and verifiable answers to complex queries based on their own datasets within a project.

- At the management layer:

The Red Hat OpenShift AI Operator

A meta-operator that deploys and maintains all components and sub-operators that are part of OpenShift AI.

When you install the Red Hat OpenShift AI Operator in the OpenShift cluster using the predefined projects, the following new projects are created:

- The **redhat-ods-operator** project contains the Red Hat OpenShift AI Operator.

- The **redhat-ods-applications** project includes the dashboard and other required components of OpenShift AI.
- The **rhods-notebooks** project is where basic workbenches are deployed by default.

You can specify custom projects if needed. You or your data scientists must also create additional projects for the applications that will use your machine learning models.

Do not install independent software vendor (ISV) applications in namespaces associated with OpenShift AI.

CHAPTER 2. UNDERSTANDING UPDATE CHANNELS

You can use update channels to specify which Red Hat OpenShift AI minor version you intend to update your Operator to. Update channels also allow you to choose the timing and level of support your updates have through the **fast**, **stable**, **stable-x.y**, **eus-x.y**, and **alpha** channel options.

The subscription of an installed Operator specifies the update channel, which is used to track and receive updates for the Operator. You can change the update channel to start tracking and receiving updates from a newer channel. For more information about the release frequency and the lifecycle associated with each of the available update channels, see the [Red Hat OpenShift AI Self-Managed Life Cycle](#) Knowledgebase article.


Channel	Support	Release frequency	Recommended environment
fast or fast-x.y	One month of full support	Every month	Production environments with access to the latest product features. Select this streaming channel with automatic updates to avoid manually upgrading every month. NOTE: OpenShift AI 3.0 is available only through the fast-3.x channel, not the general fast channel.
stable	Three months of full support	Every three months	Production environments with stability prioritized over new feature availability. Select this streaming channel with automatic updates to access the latest stable release and avoid manually upgrading.
stable-x.y	Seven months of full support	Every three months	Production environments with stability prioritized over new feature availability. Select numbered stable channels (such as stable-2.25) to plan and upgrade to the next stable release while keeping your deployment under full support.
eus-x.y	Seven months of full support followed by Extended Update Support for eleven months	Every nine months	Enterprise-grade environments that cannot upgrade within a seven month window. Select this streaming channel if you prioritize stability over new feature availability.

Channel	Support	Release frequency	Recommended environment
alpha	One month of full support	Every month	Development environments with early-access features that might not be functionally complete. Select this channel to use early-access features to test functionality and provide feedback during the development process. Early-access features are not supported with Red Hat production service level agreements (SLAs). For more information about the support scope of Red Hat Technology Preview features, see Technology Preview Features Support Scope. For more information about the support scope of Red Hat Developer Preview features, see Developer Preview Features Support Scope.

**NOTE**

The **embedded** and **beta** channels are legacy channels that will be removed in a future release. Do not select the **embedded** or **beta** channels for a new Operator installation.

CHAPTER 3. INSTALLING AND DEPLOYING OPENSIFT AI



IMPORTANT

You cannot upgrade from OpenShift AI 2.25 or any earlier version to 3.0. OpenShift AI 3.0 introduces significant technology and component changes and is intended for new installations only. To use OpenShift AI 3.0, install the Red Hat OpenShift AI Operator on a cluster running OpenShift Container Platform 4.19 or later and select the **fast-3.x** channel.

Support for upgrades will be available in a later release, including upgrades from OpenShift AI 2.25 to a stable 3.x version.

For more information, see the [Why upgrades to OpenShift AI 3.0 are not supported](#) Knowledgebase article.

Red Hat OpenShift AI is a platform for data scientists and developers of artificial intelligence (AI) applications. It provides a fully supported environment that lets you rapidly develop, train, test, and deploy machine learning models on-premises and/or in the public cloud.

OpenShift AI is provided as a managed cloud service add-on for Red Hat OpenShift or as self-managed software that you can install on-premise or in the public cloud on OpenShift.

For information about installing OpenShift AI as self-managed software on your OpenShift cluster in a disconnected environment, see [Installing and uninstalling OpenShift AI Self-Managed in a disconnected environment](#). For information about installing OpenShift AI as a managed cloud service add-on, see [Installing and uninstalling OpenShift AI Cloud Service](#).

Installing OpenShift AI involves the following high-level tasks:

1. Confirm that your OpenShift cluster meets all requirements. See [Requirements for OpenShift AI Self-Managed](#).
2. Install the Red Hat OpenShift AI Operator. See [Installing the Red Hat OpenShift AI Operator](#).
3. Install OpenShift AI components. See [Installing and managing Red Hat OpenShift AI components](#).
4. Complete any additional configuration required for the components you enabled. See the component-specific configuration sections for details.
5. Configure user and administrator groups to provide user access to OpenShift AI. See [Adding users to OpenShift AI user groups](#).
6. Access the OpenShift AI dashboard. See [Accessing the OpenShift AI dashboard](#).

3.1. REQUIREMENTS FOR OPENSIFT AI SELF-MANAGED

You must meet the following requirements before you can install Red Hat OpenShift AI on your Red Hat OpenShift cluster.

3.1.1. Platform requirements

Subscriptions

- A subscription for Red Hat OpenShift AI Self-Managed is required.
- If you want to install OpenShift AI Self-Managed in a Red Hat-managed cloud environment, you must also have a subscription for one of the following platforms:
 - Red Hat OpenShift Dedicated on Amazon Web Services (AWS) or Google Cloud Platform (GCP)
 - Red Hat OpenShift Service on Amazon Web Services (ROSA classic)
 - Red Hat OpenShift Service on Amazon Web Services with hosted control planes (ROSA HCP)
 - Microsoft Azure Red Hat OpenShift

Contact your Red Hat account manager to purchase new subscriptions. If you do not yet have an account manager, complete the form at <https://www.redhat.com/en/contact> to request one.

Cluster administrator access

- Cluster administrator access is required to install OpenShift AI.
- You can use an existing cluster or create a new one that meets the supported version requirements.

Supported OpenShift versions

The following OpenShift versions are supported for installing OpenShift AI:

- OpenShift Container Platform 4.19 to 4.20. See [OpenShift Container Platform installation overview](#).
 - To deploy models by using Distributed Inference with llm-d, your cluster must be running version 4.19.9 or later.
- OpenShift Dedicated 4. See [Creating an OpenShift Dedicated cluster](#).
- ROSA classic 4. See [Install ROSA classic clusters](#).
- ROSA HCP 4. See [Install ROSA with HCP clusters](#).



IMPORTANT

In OpenStack, CodeReady Containers (CRC), and other private cloud environments without integrated external DNS, you must manually configure DNS A or CNAME records after installing the Operator and components, when the LoadBalancer IP becomes available. For more information, see [Configuring External DNS for RHOAI 3.x on OpenStack and Private Clouds](#).

Cluster configuration

- A minimum of 2 worker nodes with at least 8 CPUs and 32 GiB RAM each is required to install the Operator.
- For single-node OpenShift clusters, the node must have at least 32 CPUs and 128 GiB RAM.
- Additional resources are required depending on your workloads.

- Open Data Hub must not be installed on the cluster.

Storage requirements

- Your cluster must have a default storage class that supports dynamic provisioning. To confirm that a default storage class is configured, run the following command:

```
oc get storageclass
```

If no storage class is marked as the default, see [Changing the default storage class](#) in the OpenShift Container Platform documentation.

Identity provider configuration

- An identity provider must be configured for your OpenShift cluster, which provides authentication for OpenShift AI. See [Understanding identity provider configuration](#).
- You must access the cluster as a user with the **cluster-admin** role; the **kubeadmin** user is not allowed. For more information, see the relevant documentation:
 - OpenShift Container Platform: [Creating a cluster admin](#)
 - OpenShift Dedicated: [Managing OpenShift Dedicated administrators](#)
 - ROSA: [Creating a cluster administrator user for quick cluster access](#)

Internet access

- Along with internet access, the following domains must be accessible during the installation of OpenShift AI:
 - **cdn.redhat.com**
 - **subscription.rhn.redhat.com**
 - **registry.access.redhat.com**
 - **registry.redhat.io**
 - **quay.io**
- For environments that build or customize CUDA-based images using NVIDIA's base images, or that directly pull artifacts from the NVIDIA NGC catalog, the following domains must also be accessible:
 - **ngc.download.nvidia.cn**
 - **developer.download.nvidia.com**



NOTE

Access to these NVIDIA domains is not required for standard OpenShift AI installations. The CUDA-based container images used by OpenShift AI are prebuilt and hosted on Red Hat's registry at **registry.redhat.io**.

Object storage

- Several components of OpenShift AI require or can use S3-compatible object storage, such as AWS S3, MinIO, Ceph, or IBM Cloud Storage. Object storage provides HTTP-based access to data by using the S3 API, which is the standard interface for most object storage services.
- Object storage is required for:
 - Single-model serving platform, for storing and deploying models.
 - AI pipelines, for storing artifacts, logs, and intermediate results.
- Object storage can also be used by:
 - Workbenches, for accessing large datasets.
 - Kueue-based workloads, for reading input data and writing output results.
 - Code executed inside pipelines, for persisting generated models or other artifacts.

Custom namespaces

- By default, OpenShift AI uses predefined namespaces, but you can define custom namespaces for the Operator, applications, and workbenches if needed. Namespaces created by OpenShift AI typically include **openshift** or **redhat** in their name. Do not rename these system namespaces because they are required for OpenShift AI to function properly.
- If you use custom namespaces, create and label them before installing the OpenShift AI Operator. See [Configuring custom namespaces](#).

3.1.2. Component requirements

Meet the requirements for the components and capabilities that you plan to use.

Workbenches (**workbenches**)

- To use a custom workbench namespace, create the namespace before installing the OpenShift AI Operator. See [Configuring custom namespaces](#).

AI Pipelines (**aipipelines**)

- To store your pipeline artifacts in an S3-compatible object storage bucket so that you do not consume local storage, configure write access to your S3 bucket on your storage account.
- If your cluster is running in FIPS mode, any custom container images for data science pipelines must be based on UBI 9 or RHEL 9. This ensures compatibility with FIPS-approved pipeline components and prevents errors related to mismatched OpenSSL or GNU C Library (glibc) versions.
- To use your own Argo Workflows instance, after installing the OpenShift AI Operator see [Configuring pipelines with your own Argo Workflows instance](#).

Kueue-based workloads (**kueue, ray, trainingoperator**)

- Install the Red Hat build of Kueue Operator.

- Install the cert-manager Operator.
- See [Configuring workload management with Kueue](#) and [Installing the distributed workloads components](#).

Model serving platform (kserve)

- Install the cert-manager Operator.

Distributed Inference with llm-d (advanced kserve)

- Install the cert-manager Operator.
- Install the Red Hat Connectivity Link Operator.
- Install the Red Hat Leader Worker Set Operator.
- See [Deploying models by using Distributed Inference with llm-d](#).

Llama Stack and RAG workloads (llamastackoperator)

- Install the Llama Stack Operator.
- Install the Red Hat OpenShift Service Mesh Operator 3.x.
- Install the cert-manager Operator.
- Ensure you have GPU-enabled nodes available on your cluster.
- Install the Node Feature Discovery Operator.
- Install the NVIDIA GPU Operator.
- Configure access to S3-compatible object storage for your model artifacts.
- See [Working with Llama Stack](#).

Model registry (modelregistry)

- Configure access to an external MySQL database 5.x or later; 8.x is recommended.
- Configure access to S3-compatible object storage.
- See [Creating a model registry](#).

3.2. CONFIGURING CUSTOM NAMESPACES

By default, OpenShift AI uses the following predefined namespaces:

- **redhat-ods-operator** contains the Red Hat OpenShift AI Operator
- **redhat-ods-applications** includes the dashboard and other required components of OpenShift AI
- **rhods-notebooks** is where basic workbenches are deployed by default

If needed, you can define custom namespaces to use instead of the predefined ones before installing OpenShift AI. This flexibility supports environments with naming policies or conventions and allows cluster administrators to control where components such as workbenches are deployed.

Namespaces created by OpenShift AI typically include **openshift** or **redhat** in their name. Do not rename these system namespaces because they are required for OpenShift AI to function properly.

Prerequisites

- You have access to an OpenShift AI cluster with cluster administrator privileges.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS
- You have not yet installed the Red Hat OpenShift AI Operator.

Procedure

1. In a terminal window, if you are not already logged in to your OpenShift cluster as a cluster administrator, log in to the OpenShift CLI (**oc**) as shown in the following example:

```
oc login <openshift_cluster_url> -u <admin_username> -p <password>
```

2. Optional: To configure a custom operator namespace:
 - a. Create a namespace YAML file named **operator-namespace.yaml**.

```
apiVersion: v1
kind: Namespace
metadata:
  name: <operator-namespace> 1
```

1 Defines the operator namespace.

- b. Create the namespace in your OpenShift cluster.

```
$ oc create -f operator-namespace.yaml
```

You see output similar to the following:

```
namespace/<operator-namespace> created
```

- c. When you install the Red Hat OpenShift AI Operator, use this namespace instead of **redhat-ods-operator**.
3. Optional: To configure a custom applications namespace:
 - a. Create a namespace YAML file named **applications-namespace.yaml**.

```
apiVersion: v1
```

```
kind: Namespace
metadata:
  name: <applications-namespace> 1
  labels:
    opendatahub.io/application-namespace: 'true' 2
```

1 Defines the applications namespace.

2 Adds the required label.

b. Create the namespace in your OpenShift cluster.

```
$ oc create -f applications-namespace.yaml
```

You see output similar to the following:

```
namespace/<applications-namespace> created
```

4. Optional: To configure a custom workbench namespace:

a. Create a namespace YAML file named **workbench-namespace.yaml**.

```
apiVersion: v1
kind: Namespace
metadata:
  name: <workbench-namespace> 1
```

1 Defines the workbench namespace.

b. Create the namespace in your OpenShift cluster.

```
$ oc create -f workbench-namespace.yaml
```

You see output similar to the following:

```
namespace/<workbench-namespace> created
```


c. When you install the Red Hat OpenShift AI components, specify this namespace for the **spec.workbenches.workbenchNamespace** field. You cannot change the default workbench namespace after you have installed the Red Hat OpenShift AI Operator.

Next step

- [Installing the Red Hat OpenShift AI Operator](#)

3.3. INSTALLING THE RED HAT OPENSIFT AI OPERATOR

This section shows how to install the Red Hat OpenShift AI Operator on your OpenShift cluster using the command-line interface (CLI) and the OpenShift web console.

**NOTE**


If your OpenShift cluster uses a proxy to access the Internet, you can configure the proxy settings for the Red Hat OpenShift AI Operator. See [Overriding proxy settings of an Operator](#) for more information.

3.3.1. Installing the Red Hat OpenShift AI Operator by using the CLI

The following procedure shows how to use the OpenShift CLI (**oc**) to install the Red Hat OpenShift AI Operator on your OpenShift cluster. You must install the Operator before you can install OpenShift AI components on the cluster.

Prerequisites

- You have a running OpenShift cluster, version 4.19 or greater, configured with a default storage class that can be dynamically provisioned.
- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS
- If you are using custom namespaces, you have created and labeled them as required.

**NOTE**

The example commands in this procedure use the predefined operator namespace. If you are using a custom operator namespace, replace **redhat-ods-operator** with your namespace.

Procedure

1. Open a new terminal window.
2. Follow these steps to log in to your OpenShift cluster as a cluster administrator:
 - a. In the upper-right corner of the OpenShift web console, click your user name and select **Copy login command**.
 - b. After you have logged in, click **Display token**.
 - c. Copy the **Log in with this token** command and paste it in your terminal.

```
$ oc login --token=<token> --server=<openshift_cluster_url>
```

3. Create a namespace for installation of the Operator by performing the following actions:

**NOTE**

If you have already created a custom namespace for the Operator, you can skip this step.

- a. Create a namespace YAML file named **rhods-operator-namespace.yaml**.

```
apiVersion: v1
kind: Namespace
metadata:
  name: redhat-ods-operator ❶
```

- ❶ Defines the operator namespace.

- b. Create the namespace in your OpenShift cluster.

```
$ oc create -f rhods-operator-namespace.yaml
```

You see output similar to the following:

```
namespace/redhat-ods-operator created
```

4. Create an operator group for installation of the Operator by performing the following actions:

- a. Create an **OperatorGroup** object custom resource (CR) file, for example, **rhods-operator-group.yaml**.

```
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: rhods-operator
  namespace: redhat-ods-operator ❶
```

- ❶ Defines the operator namespace.

- b. Create the **OperatorGroup** object in your OpenShift cluster.

```
$ oc create -f rhods-operator-group.yaml
```

You see output similar to the following:

```
operatorgroup.operators.coreos.com/rhods-operator created
```

5. Create a subscription for installation of the Operator by performing the following actions:

- a. Create a **Subscription** object CR file, for example, **rhods-operator-subscription.yaml**.

```
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: rhods-operator
  namespace: redhat-ods-operator ❶
spec:
  name: rhods-operator
  channel: <channel> ❷
```

```
source: redhat-operators
sourceNamespace: openshift-marketplace
startingCSV: rhods-operator.x.y.z 3
```

- 1** Defines the operator namespace.
- 2** Sets the update channel. You must specify a value of **fast**, **fast-x.y**, **stable**, **stable-x.y**, **eus-x.y**, or **alpha**. For more information, see [Understanding update channels](#).
- 3** Optional: Sets the operator version. If you do not specify a value, the subscription defaults to the latest operator version. For more information, see the [Red Hat OpenShift AI Self-Managed Life Cycle](#) Knowledgebase article.

- b. Create the **Subscription** object in your OpenShift cluster to install the Operator.

```
$ oc create -f rhods-operator-subscription.yaml
```

You see output similar to the following:

```
subscription.operators.coreos.com/rhods-operator created
```

Verification

- In the OpenShift web console, click **Operators → Installed Operators** and confirm that the Red Hat OpenShift AI Operator shows one of the following statuses:
 - **Installing** - installation is in progress; wait for this to change to **Succeeded**. This might take several minutes.
 - **Succeeded** - installation is successful.

Next step

- [Installing and managing Red Hat OpenShift AI components](#)

Additional resources

- [Adding users to OpenShift AI user groups](#)
- [Adding Operators to a cluster](#)

3.3.2. Installing the Red Hat OpenShift AI Operator by using the web console

The following procedure shows how to use the OpenShift web console to install the Red Hat OpenShift AI Operator on your cluster. You must install the Operator before you can install OpenShift AI components on the cluster.

Prerequisites

- You have a running OpenShift cluster, version 4.19 or greater, configured with a default storage class that can be dynamically provisioned.
- You have cluster administrator privileges for your OpenShift cluster.

- If you are using custom namespaces, you have created and labeled them as required.

Procedure

1. Log in to the OpenShift web console as a cluster administrator.
2. In the web console, click **Operators** → **OperatorHub**.
3. On the **OperatorHub** page, locate the Red Hat OpenShift AI Operator by scrolling through the available Operators or by typing *Red Hat OpenShift AI* into the **Filter by keyword** box.
4. Click the **Red Hat OpenShift AI** tile. The **Red Hat OpenShift AI** information pane opens.
5. Select a **Channel**. For information about subscription update channels, see [Understanding update channels](#).
6. Select a **Version**.
7. Click **Install**. The **Install Operator** page opens.
8. Review or change the selected channel and version as needed.
9. For **Installation mode**, note that the only available value is **All namespaces on the cluster (default)**. This installation mode makes the Operator available to all namespaces in the cluster.
10. For **Installed Namespace**, choose one of the following options:
 - To use the predefined operator namespace, select the **Operator recommended Namespace: redhat-ods-operator** option.
 - To use the custom operator namespace that you created, select the **Select a Namespace** option, and then select the namespace from the drop-down list.
11. For **Update approval**, select one of the following update strategies:
 - **Automatic**: New updates in the update channel are installed as soon as they become available.
 - **Manual**: A cluster administrator must approve any new updates before installation begins.



IMPORTANT

By default, the Red Hat OpenShift AI Operator follows a sequential update process. This means that if there are several versions between the current version and the target version, Operator Lifecycle Manager (OLM) upgrades the Operator to each of the intermediate versions before it upgrades it to the final, target version.

If you configure automatic upgrades, OLM automatically upgrades the Operator to the *latest* available version. If you configure manual upgrades, a cluster administrator must manually approve each sequential update between the current version and the final, target version.

For information about supported versions, see the [Red Hat OpenShift AI Life Cycle](#) Knowledgebase article.

12. Click **Install**.

The **Installing Operators** pane appears. When the installation finishes, a checkmark appears next to the Operator name.

Verification

- In the OpenShift web console, click **Operators** → **Installed Operators** and confirm that the Red Hat OpenShift AI Operator shows one of the following statuses:
 - **Installing** - installation is in progress; wait for this to change to **Succeeded**. This might take several minutes.
 - **Succeeded** - installation is successful.

Next step

- [Installing and managing Red Hat OpenShift AI components](#)

Additional resources

- [Adding users to OpenShift AI user groups](#)
- [Adding Operators to a cluster](#)

3.4. INSTALLING AND MANAGING RED HAT OPENSIFT AI COMPONENTS

You can use the OpenShift command-line interface (CLI) or OpenShift web console to install and manage components of Red Hat OpenShift AI on your OpenShift cluster.

3.4.1. Installing Red Hat OpenShift AI components by using the CLI

To install Red Hat OpenShift AI components by using the OpenShift CLI (**oc**), you must create and configure a **DataScienceCluster** object.



IMPORTANT

The following procedure describes how to create and configure a **DataScienceCluster** object to install Red Hat OpenShift AI components as part of a *new* installation.

For information about changing the installation status of OpenShift AI components *after* installation, see [Updating the installation status of Red Hat OpenShift AI components by using the web console](#).

Prerequisites

- The Red Hat OpenShift AI Operator is installed on your OpenShift cluster. See [Installing the Red Hat OpenShift AI Operator](#).
- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform

- [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS
- If you are using custom namespaces, you have created the namespaces.

Procedure

1. Open a new terminal window.
2. Follow these steps to log in to your OpenShift cluster as a cluster administrator:
 - a. In the upper-right corner of the OpenShift web console, click your user name and select **Copy login command**.
 - b. After you have logged in, click **Display token**.
 - c. Copy the **Log in with this token** command and paste it in your terminal.

```
$ oc login --token=<token> --server=<openshift_cluster_url>
```

3. Create a **DataScienceCluster** object custom resource (CR) file, for example, **rhods-operator-dsc.yaml**.

```
apiVersion: datasciencecluster.opendatahub.io/v2
kind: DataScienceCluster
metadata:
  name: default-dsc
spec:
  components:
    aipipelines:
      argoWorkflowsControllers:
        managementState: Removed 1
        managementState: Removed
      dashboard:
        managementState: Removed
      feastoperator:
        managementState: Removed
      kserve:
        managementState: Removed
      kueue:
        defaultClusterQueueName: default
        defaultLocalQueueName: default
        managementState: Removed
      llamastackoperator:
        managementState: Removed
      modelregistry:
        managementState: Removed
        registriesNamespace: rhoai-model-registries
      ray:
        managementState: Removed
      trainingoperator:
        managementState: Removed
      trustyai:
        managementState: Removed
```

```
workbenches:
  managementState: Removed
  workbenchNamespace: rhods-notebooks 2
```

- 1** To use your own Argo Workflows instance with the **aipipelines** component, set **argoWorkflowsControllers.managementState** to **Removed**. This allows you to integrate with a managed Argo Workflows installation already on your OpenShift cluster and avoid conflicts with the embedded controller. See *Configuring pipelines with your own Argo Workflows instance*.
 - 2** To use the predefined workbench namespace, set this value to **rhods-notebooks** or omit this line. To use a custom workbench namespace, set this value to your namespace.
4. In the **spec.components** section of the CR, for each OpenShift AI component shown, set the value of the **managementState** field to either **Managed** or **Removed**. These values are defined as follows:

Managed

The Operator actively manages the component, installs it, and tries to keep it active. The Operator will upgrade the component only if it is safe to do so.

Removed

The Operator actively manages the component but does not install it. If the component is already installed, the Operator will try to remove it.



IMPORTANT

- To learn how to fully install the KServe component, which is used by the single-model serving platform to serve large models, see [Installing the model serving platform](#).
- To learn how to install the distributed workloads components, see [Installing the distributed workloads components](#).

5. Create the **DataScienceCluster** object in your OpenShift cluster to install the specified OpenShift AI components.

```
$ oc create -f rhods-operator-dsc.yaml
```

You see output similar to the following:

```
datasciencecluster.datasciencecluster.opendatahub.io/default created
```

Verification

1. Confirm that there is at least one running pod for each component:
 - a. In the OpenShift web console, click **Workloads** → **Pods**.
 - b. In the **Project** list at the top of the page, select **redhat-ods-applications**.
 - c. In the applications namespace, confirm that there are one or more running pods for each of the OpenShift AI components that you installed.

2. Confirm the status of all installed components:

- a. In the OpenShift web console, click **Operators → Installed Operators**.
- b. Click the **Red Hat OpenShift AI Operator**.
- c. Click the **Data Science Cluster** tab.
- d. For the **DataScienceCluster** object called **default-dsc**, verify that the status is **Phase: Ready**.



NOTE

When you edit the **spec.components** section to change the installation status of a component, the **default-dsc** status also changes. During the initial installation, it might take a few minutes for the status phase to change from **Progressing** to **Ready**. You can access the OpenShift AI dashboard before the **default-dsc** status phase is **Ready**, but all components might not be ready.

- e. Click the **default-dsc** link to display the data science cluster details.
- f. Select the **YAML** tab.
- g. In the **status.installedComponents** section, confirm that the components you installed have a status value of **true**.



NOTE

If a component shows with the **component-name: {}** format in the **spec.components** section of the CR, the component is not installed.


3. In the OpenShift AI dashboard, users can view the list of the installed OpenShift AI components, their corresponding source (upstream) components, and the versions of the installed components, as described in [Viewing installed OpenShift AI components](#).

Next steps

- If you are using OpenStack, CodeReady Containers (CRC), or other private cloud environments without integrated external DNS, manually configure DNS A or CNAME records after the LoadBalancer IP becomes available. For more information, see [Configuring External DNS for RHOAI 3.x on OpenStack and Private Clouds](#).
- Complete any additional configuration required for the components you enabled. See the component-specific configuration sections for details.

3.4.2. Installing Red Hat OpenShift AI components by using the web console

To install Red Hat OpenShift AI components by using the OpenShift web console, you must create and configure a **DataScienceCluster** object.



IMPORTANT

The following procedure describes how to create and configure a **DataScienceCluster** object to install Red Hat OpenShift AI components as part of a *new* installation.

- For information about changing the installation status of OpenShift AI components *after* installation, see [Updating the installation status of Red Hat OpenShift AI components by using the web console](#).

Prerequisites

- The Red Hat OpenShift AI Operator is installed on your OpenShift cluster. See [Installing the Red Hat OpenShift AI Operator](#).
- You have cluster administrator privileges for your OpenShift cluster.
- If you are using custom namespaces, you have created the namespaces.

Procedure

1. Log in to the OpenShift web console as a cluster administrator.
2. In the web console, click **Operators** → **Installed Operators** and then click the **Red Hat OpenShift AI Operator**.
3. Click the **Data Science Cluster** tab.
4. Click **Create DataScienceCluster**.
5. For **Configure via**, select **YAML view**.

An embedded YAML editor opens showing a default custom resource (CR) for the **DataScienceCluster** object, similar to the following example:

```
apiVersion: datasciencecluster.opendatahub.io/v2
kind: DataScienceCluster
metadata:
  name: default-dsc
spec:
  components:
    aipipelines:
      argoWorkflowsControllers:
        managementState: Removed 1
        managementState: Removed
    dashboard:
      managementState: Removed
    feastoperator:
      managementState: Removed
    kserve:
      managementState: Removed
    kueue:
      defaultClusterQueueName: default
      defaultLocalQueueName: default
      managementState: Removed
    llamastackoperator:
      managementState: Removed
    modelregistry:
```

```

managementState: Removed
registriesNamespace: rhoai-model-registries
ray:
  managementState: Removed
trainingoperator:
  managementState: Removed
trustyai:
  managementState: Removed
workbenches:
  managementState: Removed
  workbenchNamespace: rhods-notebooks 2

```

- 1 To use your own Argo Workflows instance with the **aipipelines** component, set **argoWorkflowsControllers.managementState** to **Removed**. This allows you to integrate with a managed Argo Workflows installation already on your OpenShift cluster and avoid conflicts with the embedded controller. See *Configuring pipelines with your own Argo Workflows instance*.
 - 2 To use the predefined workbench namespace, set this value to **rhods-notebooks** or omit this line. To use a custom workbench namespace, set this value to your namespace.
6. In the **spec.components** section of the CR, for each OpenShift AI component shown, set the value of the **managementState** field to either **Managed** or **Removed**. These values are defined as follows:

Managed

The Operator actively manages the component, installs it, and tries to keep it active. The Operator will upgrade the component only if it is safe to do so.

Removed

The Operator actively manages the component but does not install it. If the component is already installed, the Operator will try to remove it.



IMPORTANT

- To learn how to fully install the KServe component, which is used by the single-model serving platform to serve large models, see [Installing the model serving platform](#).
- To learn how to install the distributed workloads components, see [Installing the distributed workloads components](#).

7. Click **Create**.

Verification

1. Confirm the status of all installed components:
 - a. In the OpenShift web console, click **Operators → Installed Operators**.
 - b. Click the **Red Hat OpenShift AI Operator**.
 - c. Click the **Data Science Cluster** tab.

- d. For the **DataScienceCluster** object called **default-dsc**, verify that the status is **Phase: Ready**.



NOTE

When you edit the **spec.components** section to change the installation status of a component, the **default-dsc** status also changes. During the initial installation, it might take a few minutes for the status phase to change from **Progressing** to **Ready**. You can access the OpenShift AI dashboard before the **default-dsc** status phase is **Ready**, but all components might not be ready.

- e. Click the **default-dsc** link to display the data science cluster details.
- f. Select the **YAML** tab.
- g. In the **status.installedComponents** section, confirm that the components you installed have a status value of **true**.



NOTE

If a component shows with the **component-name: {}** format in the **spec.components** section of the CR, the component is not installed.

2. Confirm that there is at least one running pod for each component:
 - a. In the OpenShift web console, click **Workloads → Pods**.
 - b. In the **Project** list at the top of the page, select **redhat-ods-applications** or your custom applications namespace.
 - c. In the applications namespace, confirm that there are one or more running pods for each of the OpenShift AI components that you installed.
3. In the OpenShift AI dashboard, users can view the list of the installed OpenShift AI components, their corresponding source (upstream) components, and the versions of the installed components, as described in [Viewing installed OpenShift AI components](#).

Next steps

- If you are using OpenStack, CodeReady Containers (CRC), or other private cloud environments without integrated external DNS, manually configure DNS A or CNAME records after the LoadBalancer IP becomes available. For more information, see [Configuring External DNS for RHOAI 3.x on OpenStack and Private Clouds](#).
- Complete any additional configuration required for the components you enabled. See the component-specific configuration sections for details.

3.4.3. Updating the installation status of Red Hat OpenShift AI components by using the web console

You can use the OpenShift web console to update the installation status of components of Red Hat OpenShift AI on your OpenShift cluster.

Prerequisites

- The Red Hat OpenShift AI Operator is [installed](#) on your OpenShift cluster.
- You have cluster administrator privileges for your OpenShift cluster.

Procedure

1. Log in to the OpenShift web console as a cluster administrator.
2. In the web console, click **Operators** → **Installed Operators** and then click the **Red Hat OpenShift AI Operator**.
3. Click the **Data Science Cluster** tab.
4. On the **DataScienceClusters** page, click the **default-dsc** object.
5. Click the **YAML** tab.
An embedded YAML editor opens showing the default custom resource (CR) for the **DataScienceCluster** object, similar to the following example:

```
apiVersion: datasciencecluster.opendatahub.io/v2
kind: DataScienceCluster
metadata:
  name: default-dsc
spec:
  components:
    aipipelines:
      argoWorkflowsControllers:
        managementState: Removed
      managementState: Removed
    dashboard:
      managementState: Removed
    feastoperator:
      managementState: Removed
    kserve:
      managementState: Removed
    kueue:
      defaultClusterQueueName: default
      defaultLocalQueueName: default
      managementState: Removed
    llamastackoperator:
      managementState: Removed
    modelregistry:
      managementState: Removed
      registriesNamespace: rhoai-model-registries
    ray:
      managementState: Removed
    trainingoperator:
      managementState: Removed
    trustyai:
      managementState: Removed
    workbenches:
      managementState: Removed
      workbenchNamespace: rhods-notebooks
```


6. In the **spec.components** section of the CR, for each OpenShift AI component shown, set the value of the **managementState** field to either **Managed** or **Removed**. These values are defined as follows:

Managed

The Operator actively manages the component, installs it, and tries to keep it active. The Operator will upgrade the component only if it is safe to do so.

Removed

The Operator actively manages the component but does not install it. If the component is already installed, the Operator will try to remove it.



IMPORTANT

- To learn how to install the KServe component, which is used by the single-model serving platform to serve large models, see [Installing the model serving platform](#).
- To learn how to install the distributed workloads feature, see [Installing the distributed workloads components](#).

7. Click **Save**.
For any components that you updated, OpenShift AI initiates a rollout that affects all pods to use the updated image.
8. If you are upgrading from OpenShift AI 2.19 or earlier, upgrade the Authorino Operator to the **stable** update channel, version 1.2.1 or later.
 - a. Update Authorino to the latest available release in the **tech-preview-v1** channel (1.1.2), if you have not done so already.
 - b. Switch to the **stable** channel:
 - i. Navigate to the Subscription settings of the Authorino Operator.
 - ii. Under **Update channel**, click on the highlighted **tech-preview-v1**.
 - iii. Change the channel to **stable**.
 - c. Select the update option for Authorino 1.2.1.

Verification

1. Confirm that there is at least one running pod for each component:
 - a. In the OpenShift web console, click **Workloads → Pods**.
 - b. In the **Project** list at the top of the page, select **redhat-ods-applications** or your custom applications namespace.
 - c. In the applications namespace, confirm that there are one or more running pods for each of the OpenShift AI components that you installed.
2. Confirm the status of all installed components:
 - a. In the OpenShift web console, click **Operators → Installed Operators**.

- b. Click the **Red Hat OpenShift AI** Operator.
- c. Click the **Data Science Cluster** tab and select the **DataScienceCluster** object called **default-dsc**.
- d. Select the **YAML** tab.
- e. In the **status.installedComponents** section, confirm that the components you installed have a status value of **true**.



NOTE

If a component shows with the **component-name: {}** format in the **spec.components** section of the CR, the component is not installed.

3. In the OpenShift AI dashboard, users can view the list of the installed OpenShift AI components, their corresponding source (upstream) components, and the versions of the installed components, as described in [Viewing installed OpenShift AI components](#).

3.4.4. Viewing installed OpenShift AI components

In the Red Hat OpenShift AI dashboard, you can view a list of the installed OpenShift AI components, their corresponding source (upstream) components, and the versions of the installed components.

Prerequisites

- OpenShift AI is installed in your OpenShift cluster.

Procedure

1. Log in to the OpenShift AI dashboard.
2. In the top navigation bar, click the help icon () and then select **About**.

Verification

The **About** page shows a list of the installed OpenShift AI components along with their corresponding upstream components and upstream component versions.

Additional resources

- [Installing and managing Red Hat OpenShift AI components](#)

CHAPTER 4. CONFIGURING PIPELINES WITH YOUR OWN ARGO WORKFLOWS INSTANCE

You can configure OpenShift AI to use an existing Argo Workflows instance instead of the embedded one included with AI pipelines. This configuration is useful if your OpenShift cluster already includes a managed Argo Workflows instance and you want to integrate it with OpenShift AI pipelines without conflicts. Disabling the embedded Argo Workflows controller allows cluster administrators to manage the lifecycles of OpenShift AI and Argo Workflows independently.



NOTE

You cannot enable both the embedded Argo Workflows instance and your own Argo Workflows instance on the same cluster.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed Red Hat OpenShift AI.

Procedure

1. Log in to the OpenShift web console as a cluster administrator.
2. In the OpenShift console, click **Operators → Installed Operators**.
3. Search for the **Red Hat OpenShift AI** Operator, and then click the Operator name to open the Operator details page.
4. Click the **Data Science Cluster** tab.
5. Click the default instance name (for example, **default-dsc**) to open the instance details page.
6. Click the **YAML** tab to show the instance specifications.
7. Disable the embedded Argo Workflows controllers that are managed by the OpenShift AI Operator:
 - a. In the **spec.components** section, set the value of the **managementState** field for the **aipipelines** component to **Managed**.
 - b. In the **spec.components.aipipelines** section, set the value of the **managementState** field for **argoWorkflowsControllers** to **Removed**, as shown in the following example:

Example aipipelines specification

```
# ...
spec:
  components:
    aipipelines:
      argoWorkflowsControllers:
        managementState: Removed
        managementState: Managed
# ...
```

8. Click **Save** to apply your changes.
9. Install and configure a compatible version of Argo Workflows on your cluster. For compatible version information, see [Supported Configurations for 3.x](#). For installation information, see the [Argo Workflows Installation documentation](#).

Verification

1. On the **Details** tab of the **DataScienceCluster** instance (for example, **default-dsc**), verify that **AIPIipelinesReady** has a **Status** of **True**.
2. Verify that the **ds-pipeline-workflow-controller** pod does not exist:
 - a. Go to **Workloads → Pods**.
 - b. Search for the **ds-pipeline-workflow-controller** pod.
 - c. Verify that this pod does not exist. The absence of this pod confirms that the embedded Argo Workflows controller is disabled.

CHAPTER 5. INSTALLING THE DISTRIBUTED WORKLOADS COMPONENTS

To use the distributed workloads feature in OpenShift AI, you must install several components.

Prerequisites

- You have logged in to OpenShift with the **cluster-admin** role and you can access the data science cluster.
- You have installed Red Hat OpenShift AI.
- You have installed the Red Hat build of Kueue Operator on your OpenShift cluster, as described in the [Red Hat build of Kueue documentation](#).
- You have sufficient resources. In addition to the minimum OpenShift AI resources described in [Installing and deploying OpenShift AI](#) (for disconnected environments, see [Deploying OpenShift AI in a disconnected environment](#)), you need 1.6 vCPU and 2 GiB memory to deploy the distributed workloads infrastructure.
- You have installed the cert-manager Operator in OpenShift by using the web console as described in [Installing the cert-manager Operator for Red Hat OpenShift](#).
- If you want to use graphics processing units (GPUs), you have enabled GPU support in OpenShift AI. If you use NVIDIA GPUs, see [Enabling NVIDIA GPUs](#). If you use AMD GPUs, see [AMD GPU integration](#).



NOTE

In OpenShift AI, Red Hat supports the use of accelerators within the same cluster only.

Starting from Red Hat OpenShift AI 2.19, Red Hat supports remote direct memory access (RDMA) for NVIDIA GPUs only, enabling them to communicate directly with each other by using NVIDIA GPUDirect RDMA across either Ethernet or InfiniBand networks.

- If you want to use self-signed certificates, you have added them to a central Certificate Authority (CA) bundle as described in [Working with certificates](#) (for disconnected environments, see [Working with certificates](#)). No additional configuration is necessary to use those certificates with distributed workloads. The centrally configured self-signed certificates are automatically available in the workload pods at the following mount points:

- Cluster-wide CA bundle:

```
/etc/pki/tls/certs/odh-trusted-ca-bundle.crt
/etc/ssl/certs/odh-trusted-ca-bundle.crt
```

- Custom CA bundle:

```
/etc/pki/tls/certs/odh-ca-bundle.crt
/etc/ssl/certs/odh-ca-bundle.crt
```

Procedure

1. In the OpenShift console, click **Operators → Installed Operators**.
2. Search for the **Red Hat OpenShift AI** Operator, and then click the Operator name to open the Operator details page.
3. Click the **Data Science Cluster** tab.
4. Click the default instance name (for example, **default-dsc**) to open the instance details page.
5. Click the **YAML** tab to show the instance specifications.
6. Enable the required distributed workloads components. In the **spec.components** section, set the **managementState** field correctly for the required components:
 - Set **kueue** to **Unmanaged** to allow the Red Hat build of Kueue Operator to manage Kueue.
 - If you want to use the Ray framework to tune models, set **ray** to **Managed**.
 - If you want to use the Kubeflow Training Operator to tune models, set **trainingoperator** to **Managed**.
 - The list of required components depends on whether the distributed workload is run from a pipeline or workbench or both, as shown in the following table.

Table 5.1. Components required for distributed workloads

Component	Pipelines only	Workbenches only	Pipelines and workbenches
dashboard	Managed	Managed	Managed
aipipelines	Managed	Removed	Managed
kueue	Unmanaged	Unmanaged	Unmanaged
ray	Managed	Managed	Managed
trainingoperator	Managed	Managed	Managed
workbenches	Removed	Managed	Managed

7. Click **Save**. After a short time, the components with a **Managed** state are ready.

Verification

Check the status of the **kubeflow-training-operator**, **kuberay-operator**, **kueue-controller-manager**, and **openshift-kueue-operator** pods, as follows:

1. In the OpenShift console, click **Workloads → Deployments**.
2. In the **Search by name** field, enter the following search strings:

- In the **redhat-ods-applications** project, search for **kubeflow-training-operator** and **kuberay-operator**.
 - In the **openshift-kueue-operator** project, search for **kueue-controller-manager** and **openshift-kueue-operator**.
3. In each case, check the status as follows:
- a. Click the deployment name to open the deployment details page.
 - b. Click the **Pods** tab.
 - c. Check the pod status.
When the status of the pods is **Running**, the pods are ready to use.
 - d. To see more information about each pod, click the pod name to open the pod details page, and then click the **Logs** tab.

Next Step

Configure the distributed workloads feature as described in [Managing distributed workloads](#).

CHAPTER 6. ACCESSING THE DASHBOARD

After you have installed OpenShift AI and added users, you can access the URL for your OpenShift AI console and share the URL with the users to let them log in and work on their models.

Prerequisites

- You have installed OpenShift AI on your OpenShift cluster.
- You have added at least one user to the user group for OpenShift AI.

Procedure

1. Log in to OpenShift web console.
2. Click the application launcher ().
3. Right-click **Red Hat OpenShift AI** and copy the URL for your OpenShift AI instance.
4. Provide this instance URL to your data scientists to let them log in to OpenShift AI.

Verification

- Confirm that you and your users can log in to OpenShift AI by using the instance URL.

Note: In the Red Hat OpenShift AI dashboard, users can view the list of the installed OpenShift AI components, their corresponding source (upstream) components, and the versions of the installed components, as described in [Viewing installed components](#).

Additional resources

- [Logging in to OpenShift AI](#)
- [Adding users to OpenShift AI user groups](#)

CHAPTER 7. ENABLING ACCELERATORS

Before you can use an accelerator in OpenShift AI, you must install the relevant software components. The installation process varies based on the accelerator type.

Prerequisites

- You have logged in to your OpenShift cluster.
- You have the **cluster-admin** role in your OpenShift cluster.
- You have installed an accelerator and confirmed that it is detected in your environment.

Procedure

1. Follow the appropriate documentation to enable your accelerator:
 - **NVIDIA GPUs:** See [Enabling NVIDIA GPUs](#).
 - **Intel Gaudi AI accelerators:** See [Enabling Intel Gaudi AI accelerators](#).
 - **AMD GPUs:** See [Enabling AMD GPUs](#).
2. After installing your accelerator, create a hardware profile as described in: [Working with hardware profiles](#).

Verification

- From the **Administrator** perspective, go to the **Operators → Installed Operators** page. Confirm that the following Operators appear:
 - The Operator for your accelerator
 - Node Feature Discovery (NFD)
 - Kernel Module Management (KMM)
- The accelerator is correctly detected a few minutes after full installation of the Node Feature Discovery (NFD) and the relevant accelerator Operator. The OpenShift CLI (**oc**) displays the appropriate output for the GPU worker node. For example, here is output confirming that an NVIDIA GPU is detected:

```
# Expected output when the accelerator is detected correctly
oc describe node <node name>
...
Capacity:
  cpu:                4
  ephemeral-storage:  313981932Ki
  hugepages-1Gi:      0
  hugepages-2Mi:      0
  memory:              16076568Ki
  nvidia.com/gpu:      1
  pods:                250
Allocatable:
  cpu:                 3920m
  ephemeral-storage:  288292006229
```



```
hugepages-1Gi:    0
hugepages-2Mi:    0
memory:           12828440Ki
nvidia.com/gpu:    1
pods:              250
```

CHAPTER 8. WORKING WITH CERTIFICATES

When you install Red Hat OpenShift AI, OpenShift automatically applies a default Certificate Authority (CA) bundle to manage authentication for most OpenShift AI components, such as workbenches and model servers. These certificates are trusted self-signed certificates that help secure communication. However, as a cluster administrator, you might need to configure additional self-signed certificates to use some components, such as the AI pipeline server and object storage solutions. If an OpenShift AI component uses a self-signed certificate that is not part of the existing cluster-wide CA bundle, you have the following options for including the certificate:

- Add it to the OpenShift cluster-wide CA bundle.
- Add it to a custom CA bundle, separate from the cluster-wide CA bundle.

As a cluster administrator, you can also change how to manage authentication for OpenShift AI as follows:

- Manually manage certificate changes, instead of relying on the OpenShift AI Operator to handle them automatically.
- Remove the cluster-wide CA bundle, either from all namespaces or specific ones. If you prefer to implement a different authentication approach, you can override the default OpenShift AI behavior, as described in *Removing the CA bundle*.

8.1. UNDERSTANDING HOW OPENSIFT AI HANDLES CERTIFICATES

After installing OpenShift AI, the Red Hat OpenShift AI Operator automatically creates an empty **odh-trusted-ca-bundle** configuration file (ConfigMap). The Cluster Network Operator (CNO) injects the cluster-wide CA bundle into the **odh-trusted-ca-bundle** configMap with the label **"config.openshift.io/inject-trusted-cabundle"**.

```
apiVersion: v1
kind: ConfigMap
metadata:
  labels:
    app.kubernetes.io/part-of: opendatahub-operator
    config.openshift.io/inject-trusted-cabundle: 'true'
name: odh-trusted-ca-bundle
```

After the CNO operator injects the bundle, it updates the ConfigMap with the contents of the **ca-bundle.crt** file.

```
apiVersion: v1
kind: ConfigMap
metadata:
  labels:
    app.kubernetes.io/part-of: opendatahub-operator
    config.openshift.io/inject-trusted-cabundle: 'true'
name: odh-trusted-ca-bundle
data:
  ca-bundle.crt: |
    <BUNDLE OF CLUSTER-WIDE CERTIFICATES>
```

The management of CA bundles is configured through the Data Science Cluster Initialization (DSCI) object. Within this object, you can set the **spec.trustedCABundle.managementState** field to one of the following values:

- **Managed:** (Default) The Red Hat OpenShift AI Operator manages the **odh-trusted-ca-bundle** ConfigMap and adds it to all non-reserved existing and new namespaces. It does not add the ConfigMap to any reserved or system namespaces, such as **default**, **openshift-*** or **kube-***. The Red Hat OpenShift AI Operator automatically updates the ConfigMap to reflect any changes made to the **customCABundle** field.
- **Unmanaged:** The Red Hat OpenShift AI administrator manually manages the **odh-trusted-ca-bundle** ConfigMap, instead of allowing the Operator to manage it. Changing the **managementState** from **Managed** to **Unmanaged** does not remove the **odh-trusted-ca-bundle** ConfigMap. However, the ConfigMap is no longer automatically updated if changes are made to the **customCABundle** field.

The **Unmanaged** setting is useful if your organization implements a different method for managing trusted CA bundles, such as Ansible automation, and does not want the Red Hat OpenShift AI Operator to handle certificates automatically. This setting provides greater control, preventing the Operator from overwriting custom configurations.

- **Removed:** The Red Hat OpenShift AI Operator removes the **odh-trusted-ca-bundle** ConfigMap, if present, and prevents ConfigMaps from being created in new namespaces. Changing this field from **Managed** to **Removed** also deletes the ConfigMap from existing namespaces. This is the default value after upgrading Red Hat OpenShift AI from 2.7 or earlier versions to 3.0.

The **Removed** setting reduces complexity and mitigates security risks, such as unauthorized certificate changes. In high-security environments, removing the CA bundle ensures that only approved CAs are trusted, reducing the risk of cyberattacks. For example, your organization might want to restrict cluster administrators from creating trusted CA bundles to prevent OpenShift pods from communicating externally.

8.2. ADDING CERTIFICATES

If you must use a self-signed certificate that is not part of the existing cluster-wide CA bundle, you have two options for configuring the certificate:

- Add it to the cluster-wide CA bundle.
This option is useful when the certificate is needed for secure communication across multiple services or when it's required by security policies to be trusted cluster-wide. This option ensures that all services and components in the cluster trust the certificate automatically. It simplifies management because the certificate is trusted across the entire cluster, avoiding the need to configure the certificate separately for each service.
- Add it to a custom CA bundle that is separate from the OpenShift cluster-wide bundle. Consider this option for the following scenarios:
 - Limit scope: Only specific services need the certificate, not the whole cluster.
 - Isolation: Keeps custom certificates separate, preventing changes to the global configuration.
 - Avoid global impact: Does not affect services that do not need the certificate.
 - Easier management: Makes it simpler to manage certificates for specific services.

8.3. ADDING CERTIFICATES TO A CLUSTER-WIDE CA BUNDLE

You can add a self-signed certificate to a cluster-wide Certificate Authority (CA) bundle (**ca-bundle.crt**).

When the cluster-wide CA bundle is updated, the Cluster Network Operator (CNO) automatically detects the change and injects the updated bundle into the **odh-trusted-ca-bundle** ConfigMap, making the certificate available to OpenShift AI components.

Note: By default, the management state for the Trusted CA bundle is **Managed** (that is, the **spec.trustedCABundle.managementState** field in the Red Hat OpenShift AI Operator's DSCI object is set to **Managed**). If you change this setting to **Unmanaged**, you must manually update the **odh-trusted-ca-bundle** ConfigMap to include the updated cluster-wide CA bundle.

Alternatively, you can add certificates to a custom CA bundle, as described in [Adding certificates to a custom CA bundle](#).

Prerequisites

- You have created a self-signed certificate and saved the certificate to a file. For example, you have created a certificate using OpenSSL and saved it to a file named **example-ca.crt**.
- You have cluster administrator access for the OpenShift cluster where Red Hat OpenShift AI is installed.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

1. Create a ConfigMap that includes the root CA certificate used to sign the certificate, where **</path/to/example-ca.crt>** is the path to the CA certificate bundle on your local file system:

```
oc create configmap custom-ca \
  --from-file=ca-bundle.crt=</path/to/example-ca.crt> \
  -n openshift-config
```

2. Update the cluster-wide proxy configuration with the newly-created ConfigMap:

```
oc patch proxy/cluster \
  --type=merge \
  --patch='{ "spec": { "trustedCA": { "name": "custom-ca" } } }'
```

Verification

Run the following command to verify that all non-reserved namespaces contain the **odh-trusted-ca-bundle** ConfigMap:

```
oc get configmaps --all-namespaces -l app.kubernetes.io/part-of=opendatahub-operator | grep odh-trusted-ca-bundle
```

Additional resources

- [Configuring certificates](#) in the OpenShift Container Platform documentation
- [Injecting a custom CA Bundle](#) in the Red Hat OpenShift Service on AWS documentation
- [Injecting a custom CA Bundle](#) in the OpenShift Dedicated documentation

8.4. ADDING CERTIFICATES TO A CUSTOM CA BUNDLE

You can add self-signed certificates to a custom CA bundle that is separate from the OpenShift cluster-wide bundle.

This method is ideal for scenarios where components need access to external resources that require a self-signed certificate. For example, you might need to add self-signed certificates to grant AI pipelines access to S3-compatible object storage.

Prerequisites

- You have created a self-signed certificate and saved the certificate to a file. For example, you have created a certificate using OpenSSL and saved it to a file named **example-ca.crt**.
- You have cluster administrator access for the OpenShift cluster where Red Hat OpenShift AI is installed.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

1. Log in to OpenShift.
2. Click **Operators** → **Installed Operators** and then click the Red Hat OpenShift AI Operator.
3. Click the **DSC Initialization** tab.
4. Click the **default-dsci** object.
5. Click the **YAML** tab.
6. In the **spec.trustedCABundle** section, add the custom certificate to the **customCABundle** field, as shown in the following example:

```
spec:
  trustedCABundle:
    managementState: Managed
    customCABundle: |
      -----BEGIN CERTIFICATE-----
      examplebundle123
      -----END CERTIFICATE-----
```

7. Click **Save**.

The Red Hat OpenShift AI Operator automatically updates the ConfigMap to reflect any changes made to the **customCABundle** field. It adds the **odh-ca-bundle.crt** file containing the certificates to the **odh-trusted-ca-bundle** ConfigMap, as shown in the following example:

```
apiVersion: v1
kind: ConfigMap
metadata:
  labels:
    app.kubernetes.io/part-of: opendatahub-operator
    config.openshift.io/inject-trusted-cabundle: 'true'
  name: odh-trusted-ca-bundle
data:
  ca-bundle.crt: |
    <BUNDLE OF CLUSTER-WIDE CERTIFICATES>
  odh-ca-bundle.crt: |
    <BUNDLE OF CUSTOM CERTIFICATES>
```

Verification

Run the following command to verify that a non-reserved namespace contains the **odh-trusted-ca-bundle** ConfigMap and that the ConfigMap contains your **customCABundle** value. In the following command, *example-namespace* is the non-reserved namespace and *examplebundle123* is the **customCABundle** value.

```
oc get configmap odh-trusted-ca-bundle -n example-namespace -o yaml | grep examplebundle123
```

8.5. USING SELF-SIGNED CERTIFICATES WITH OPENSIFT AI COMPONENTS

Some OpenShift AI components have additional options or required configuration for self-signed certificates.

8.5.1. Accessing S3-compatible object storage with self-signed certificates

To securely connect OpenShift AI components to object storage solutions or databases that are deployed within an OpenShift cluster that uses self-signed certificates, you must provide a certificate authority (CA) certificate. Each namespace includes a ConfigMap named **kube-root-ca.crt**, which contains the CA certificate of the internal API Server.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS
- You have deployed an object storage solution or database in your OpenShift cluster.

Procedure

1. In a terminal window, log in to the OpenShift CLI (**oc**) as shown in the following example:

```
oc login api.<cluster_name>.<cluster_domain>:6443 --web
```

2. Retrieve the current OpenShift AI trusted CA configuration and store it in a new file:

```
oc get dscinitializations.dscinitialization.opendatahub.io default-dsci -o json | jq -r  
'spec.trustedCABundle.customCABundle' > /tmp/my-custom-ca-bundles.crt
```

3. Add the cluster's **kube-root-ca.crt** ConfigMap to the OpenShift AI trusted CA configuration:

```
oc get configmap kube-root-ca.crt -o jsonpath="{['data']['ca.crt']}" >> /tmp/my-custom-ca-bundles.crt
```

4. Update the OpenShift AI trusted CA configuration to trust certificates issued by the certificate authorities in **kube-root-ca.crt**:

```
oc patch dscinitialization default-dsci --type='json' -  
p='[{"op":"replace","path":"/spec/trustedCABundle/customCABundle","value":""$(awk '{printf  
"%s\\n", $0}' /tmp/my-custom-ca-bundles.crt)"}]'
```

Verification

- You can successfully deploy components that are configured to use object storage solutions or databases that are deployed in the OpenShift cluster. For example, a pipeline server that is configured to use a database deployed in the cluster starts successfully.



NOTE

You can verify your new certificate configuration by following the steps in the *OpenShift AI tutorial – Fraud Detection example*. Run the script to install local object storage buckets and create connections, and then enable AI pipelines.

For more information about running the script to install local object storage buckets, see [Running a script to install local object storage buckets and create connections](#).

For more information about enabling AI pipelines, see [Enabling pipelines](#).

8.5.2. Configuring a certificate for pipelines

By default, OpenShift AI includes OpenShift cluster-wide certificates in the **odh-trusted-ca-bundle** ConfigMap. These cluster-wide certificates cover most components, such as workbenches and model servers. However, the pipeline server might require additional Certificate Authority (CA) configuration, especially when interacting with external systems that use self-signed or custom certificates.

You have the following options for adding the certificate for AI pipelines:

- Add them to the cluster-wide CA bundle, as described in [Adding certificates to a cluster-wide CA bundle](#).
- Add them to a custom bundle as described in [Adding certificates to a custom CA bundle](#).
- Provide a CA bundle that is only used for AI pipelines, as described in the following procedure.

Prerequisites

- You have cluster administrator access for the OpenShift cluster where Red Hat OpenShift AI is installed.
- You have created a self-signed certificate and saved the certificate to a file. For example, you have created a certificate using OpenSSL and saved it to a file named **example-ca.crt**.
- You have configured an AI pipeline server.

Procedure

1. Log in to the OpenShift console.
2. From **Workloads** → **ConfigMaps**, create a ConfigMap with the required bundle in the same project as the target AI pipeline:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: custom-ca-bundle
data:
  ca-bundle.crt: |
    # contents of ca-bundle.crt
```

3. Add the following snippet to the **.spec.apiserver.caBundle** field of the underlying **DataSciencePipelinesApplication** (DSPA):

```
apiVersion: datasciencepipelinesapplications.opendatahub.io/v1
kind: DataSciencePipelinesApplication
metadata:
  name: data-science-dspa
spec:
  ...
  apiServer:
  ...
  cABundle:
    configMapName: custom-ca-bundle
    configMapKey: ca-bundle.crt
```

4. Save the ConfigMap. The pipeline server pod automatically redeploys with the updated bundle.


Verification

Confirm that your CA bundle was successfully mounted:

1. Log in to the OpenShift console.
2. Go to the project that has the target AI pipeline.
3. Click the **Pods** tab.
4. Click the pipeline server pod with the **ds-pipeline-dspa-**<hash>**** prefix.
5. Click **Terminal**.

6. Enter **cat /dsp-custom-certs/dsp-ca.crt**.
7. Verify that your CA bundle is present within this file.

8.5.3. Configuring a certificate for workbenches



IMPORTANT

By default, self-signed certificates apply to workbenches that you create after configuring cluster-wide certificates. To apply cluster-wide certificates to an existing workbench, stop and then restart the workbench.

Self-signed certificates are stored in **/etc/pki/tls/custom-certs/ca-bundle.crt**. Workbenches use a preset environment variable that many popular HTTP client packages point to for certificates. For packages that are not included by default, you can provide this certificate path. For example, for the **kfp** package to connect to the AI pipeline server:

```
from kfp.client import Client

with open(sa_token_file_path, 'r') as token_file:
    bearer_token = token_file.read()

client = Client(
    host='https://<GO_TO_ROUTER_OF_PROJECT>',
    existing_token=bearer_token,
    ssl_ca_cert='/etc/pki/tls/custom-certs/ca-bundle.crt'
)
print(client.list_experiments())
```

8.5.4. Using the cluster-wide CA bundle for the model serving platform

By default, the model serving platform in OpenShift AI uses a self-signed certificate generated at installation for the endpoints that are created when deploying a server.

If you have configured cluster-wide certificates on your OpenShift cluster, they are used by default for other types of endpoints, such as endpoints for routes.

The following procedure explains how to use the same certificate that you already have for your OpenShift cluster.

Prerequisites

- You have cluster administrator access for the OpenShift cluster where Red Hat OpenShift AI is installed.
- You have configured cluster-wide certificates in OpenShift.
- You have configured the model serving platform, as described in [Installing the model serving platform](#).

Procedure

1. Log in to the OpenShift console.

2. From the list of projects, open the **openshift-ingress** project.
3. Click **YAML**.
4. Search for "cert" to find a secret with a name that includes "cert". For example, **rhods-internal-primary-cert-bundle-secret**. The contents of the secret should contain two items that are used for all OpenShift Routes: **tls.cert** (the certificate) and **tls.key** (the key).
5. Copy the reference to the secret.
6. From the list of projects, open the **istio-system** project.
7. Create a YAML file and paste the reference to the secret that you copied from the **openshift-ingress** YAML file.
8. Edit the YAML code to keep only the relevant content, as shown in the following example. Replace **rhods-internal-primary-cert-bundle-secret** with the name of your secret:

```
kind: Secret
apiVersion: v1
metadata:
  name: rhods-internal-primary-cert-bundle-secret
data:
  tls.crt: >-
    LS0tLS1CRUd...
  tls.key: >-
    LS0tLS1CRUd...
type: kubernetes.io/tls
```

9. Save the YAML file in the **istio-system** project.
10. Navigate to **Operators → Installed Operators → Red Hat OpenShift AI**.
11. Click **Data Science Cluster***, and then click **default-dsc → YAML**.
12. Edit the **kserve** configuration section to refer to your secret as shown in the following example. Replace **rhods-internal-primary-cert-bundle-secret** with the name of the secret that you created in Step 8.

```
kserve:
  devFlags: {}
  managementState: Managed
  serving:
    ingressGateway:
      certificate:
        secretName: rhods-internal-primary-cert-bundle-secret
        type: Provided
      managementState: Managed
      name: knative-serving
```

8.6. MANAGING CERTIFICATES WITHOUT THE RED HAT OPENSIFT AI OPERATOR

By default, the Red Hat OpenShift AI Operator manages the **odh-trusted-ca-bundle** ConfigMap, which contains the trusted CA bundle and is applied to all non-reserved namespaces in the cluster. The Operator automatically updates this ConfigMap whenever changes are made to the CA bundle.

If your organization prefers to manage trusted CA bundles independently, for example, by using Ansible automation, you can disable this default behavior to prevent automatic updates by the Red Hat OpenShift AI Operator.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.

Procedure

1. In the OpenShift web console, click **Operators** → **Installed Operators** and then click the **Red Hat OpenShift AI Operator**.
2. Click the **DSC Initialization** tab.
3. Click the **default-dsci** object.
4. Click the **YAML** tab.
5. In the **spec** section, change the value of the **managementState** field for **trustedCABundle** to **Unmanaged**, as shown:

```
spec:
  trustedCABundle:
    managementState: Unmanaged
```

6. Click **Save**.
Changing the **managementState** from **Managed** to **Unmanaged** prevents automatic updates when the **customCABundle** field is modified, but does not remove the **odh-trusted-ca-bundle** ConfigMap.

Verification

1. In the **spec** section, set the value of the **customCABundle** field for **trustedCABundle**, for example:

```
spec:
  trustedCABundle:
    managementState: Unmanaged
    customCABundle: example123
```

2. Click **Save**.
3. Click **Workloads** → **ConfigMaps**.
4. Select a project from the project list.
5. Click the **odh-trusted-ca-bundle** ConfigMap.
6. Click the **YAML** tab and verify that the value of the **customCABundle** field did not update.

8.7. REMOVING THE CA BUNDLE

If you prefer to implement a different authentication approach for your OpenShift AI installation, you can override the default behavior by removing the CA bundle.

You have two options for removing the CA bundle:

- Remove the CA bundle from all non-reserved projects in OpenShift AI.
- Remove the CA bundle from a specific project.

8.7.1. Removing the CA bundle from all namespaces

You can remove a Certificate Authority (CA) bundle from all non-reserved namespaces in OpenShift AI. This process changes the default configuration and disables the creation of the **odh-trusted-ca-bundle** configuration file (ConfigMap), as described in [Working with certificates](#) (OpenShift AI Self-Managed) or [Working with certificates](#) (OpenShift AI Self-Managed in a disconnected environment).



NOTE

The **odh-trusted-ca-bundle** ConfigMaps are only deleted from namespaces when you set the **managementState** of **trustedCABundle** to **Removed**; deleting the DSC Initialization does not delete the ConfigMaps.

To remove a CA bundle from a single namespace only, see [Removing the CA bundle from a single namespace](#) (OpenShift AI Self-Managed) or [Removing the CA bundle from a single namespace](#) (OpenShift AI Self-Managed in a disconnected environment).

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

1. In the OpenShift web console, click **Operators** → **Installed Operators** and then click the Red Hat OpenShift AI Operator.
2. Click the **DSC Initialization** tab.
3. Click the **default-dsci** object.
4. Click the **YAML** tab.
5. In the **spec** section, change the value of the **managementState** field for **trustedCABundle** to **Removed**:

```
spec:
  trustedCABundle:
    managementState: Removed
```

6. Click **Save**.

Verification

- Run the following command to verify that the **odh-trusted-ca-bundle** ConfigMap has been removed from all namespaces:

```
oc get configmaps --all-namespaces | grep odh-trusted-ca-bundle
```

The command should not return any ConfigMaps.

8.7.2. Removing the CA bundle from a single namespace

You can remove a custom Certificate Authority (CA) bundle from individual namespaces in OpenShift AI. This process disables the creation of the **odh-trusted-ca-bundle** configuration file (ConfigMap) for the specified namespace only.

To remove a CA bundle from all namespaces, [Removing the CA bundle from all namespaces](#) (OpenShift AI Self-Managed) or [Removing the CA bundle from all namespaces](#) (OpenShift AI Self-Managed in a disconnected environment).

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

- Run the following command to remove a CA bundle from a namespace. In the following command, *example-namespace* is the non-reserved namespace.

```
oc annotate ns example-namespace security.opendatahub.io/inject-trusted-ca-bundle=false
```

Verification

- Run the following command to verify that the CA bundle has been removed from the namespace. In the following command, *example-namespace* is the non-reserved namespace.

```
oc get configmap odh-trusted-ca-bundle -n example-namespace
```

The command should return **configmaps "odh-trusted-ca-bundle" not found**.

CHAPTER 9. VIEWING LOGS AND AUDIT RECORDS

As a cluster administrator, you can use the OpenShift AI Operator logger to monitor and troubleshoot issues. You can also use OpenShift audit records to review a history of changes made to the OpenShift AI Operator configuration.

9.1. CONFIGURING THE OPENSIFT AI OPERATOR LOGGER

You can change the log level for OpenShift AI Operator components by setting the **.spec.devFlags.logmode** flag for the **DSC Initialization/DSCI** custom resource during runtime. If you do not set a **logmode** value, the logger uses the INFO log level by default.

The log level that you set with **.spec.devFlags.logmode** applies to all components, not just those in a **Managed** state.

The following table shows the available log levels:

Log level	Stacktrace level	Verbosity	Output	Timestamp type
devel or development	WARN	INFO	Console	Epoch timestamps
"" (or no logmode value set)	ERROR	INFO	JSON	Human-readable timestamps
prod or production	ERROR	INFO	JSON	Human-readable timestamps

Logs that are set to **devel** or **development** generate in a plain text console format. Logs that are set to **prod**, **production**, or which do not have a level set generate in a JSON format.

Prerequisites

- You have administrator access to the **DSCInitialization** resources in the OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

1. Log in to the OpenShift as a cluster administrator.
2. Click **Operators** → **Installed Operators** and then click the Red Hat OpenShift AI Operator.
3. Click the **DSC Initialization** tab.
4. Click the **default-dsci** object.

5. Click the **YAML** tab.
6. In the **spec** section, update the **.spec.devFlags.logmode** flag with the log level that you want to set.

```
apiVersion: dscinitialization.opendatahub.io/v1
kind: DSCInitialization
metadata:
  name: default-dsci
spec:
  devFlags:
    logmode: development
```

7. Click **Save**.

You can also configure the log level from the OpenShift CLI (**oc**) by using the following command with the **logmode** value set to the log level that you want.

```
oc patch dsci default-dsci -p '{"spec":{"devFlags":{"logmode":"development"}}}' --type=merge
```

Verification

- If you set the component log level to **devel** or **development**, logs generate more frequently and include logs at **WARN** level and above.
- If you set the component log level to **prod** or **production**, or do not set a log level, logs generate less frequently and include logs at **ERROR** level or above.

9.1.1. Viewing the OpenShift AI Operator logs

1. Log in to the OpenShift CLI (**oc**).
2. Run the following command to stream logs from all Operator pods:

```
for pod in $(oc get pods -l name=rhods-operator -n redhat-ods-operator -o name); do
  oc logs -f "$pod" -n redhat-ods-operator &
done
```

The Operator pod logs open in your terminal.

TIP

Press **Ctrl+C** to stop viewing. To fully stop all log streams, run **kill \$(jobs -p)**.

You can also view each Operator pod log in the OpenShift console by navigating to **Workloads → Pods**, selecting the **redhat-ods-operator** project, clicking a pod name, and then clicking the **Logs** tab.

9.2. VIEWING AUDIT RECORDS

Cluster administrators can use OpenShift auditing to see changes made to the OpenShift AI Operator configuration by reviewing modifications to the DataScienceCluster (DSC) and DSCInitialization (DSCI) custom resources. Audit logging is enabled by default in standard OpenShift cluster configurations. For

more information, see [Viewing audit logs](#) in the OpenShift documentation.



NOTE

In Red Hat OpenShift Service on AWS, audit logging is disabled by default because the Elasticsearch log store does not provide secure storage for audit logs. To configure log forwarding, see [Logging](#) in the Red Hat OpenShift Service on AWS documentation.

The following example shows how to use the OpenShift audit logs to see the history of changes made (by users) to the DSC and DSCI custom resources.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS

Procedure

1. In a terminal window, if you are not already logged in to your OpenShift cluster as a cluster administrator, log in to the OpenShift CLI as shown in the following example:

```
$ oc login <openshift_cluster_url> -u <admin_username> -p <password>
```

2. To access the full content of the changed custom resources, set the OpenShift audit log policy to **WriteRequestBodies** or a more comprehensive profile. For more information, see [Configuring the audit log policy](#).
3. Fetch the audit log files that are available for the relevant control plane nodes. For example:

```
oc adm node-logs --role=master --path=kube-apiserver/ \
| awk '{ print $1 }' | sort -u \
| while read node ; do
  oc adm node-logs $node --path=kube-apiserver/audit.log < /dev/null
done \
| grep opendatahub > /tmp/kube-apiserver-audit-opendatahub.log
```

4. Search the files for the DSC and DSCI custom resources. For example:

```
jq 'select((.objectRef.apiGroup == "dscinitialization.opendatahub.io"
or .objectRef.apiGroup == "datasciencecluster.opendatahub.io")
and .user.username != "system:serviceaccount:redhat-ods-operator:redhat-ods-
operator-controller-manager"
and .verb != "get" and .verb != "watch" and .verb != "list")' < /tmp/kube-apiserver-
audit-opendatahub.log
```

Verification

- The commands return relevant log entries.

TIP

To configure the log retention time, see the [Logging](#) section in the OpenShift documentation.

CHAPTER 10. TROUBLESHOOTING COMMON INSTALLATION PROBLEMS

If you are experiencing difficulties installing the Red Hat OpenShift AI Operator, read this section to understand what could be causing the problem and how to resolve it.

If the problem is not included here or in the release notes, contact [Red Hat Support](#). When opening a support case, it is helpful to include debugging information about your cluster. You can collect this information by using the **must-gather** tool as described in [Must-Gather for Red Hat OpenShift AI](#) and [Gathering data about your cluster](#).

You can also adjust the log level of OpenShift AI Operator components to increase or reduce log verbosity to suit your use case. For more information, see [Configuring the OpenShift AI Operator logger](#).

10.1. THE RED HAT OPENSIFT AI OPERATOR CANNOT BE RETRIEVED FROM THE IMAGE REGISTRY

Problem

When attempting to retrieve the Red Hat OpenShift AI Operator from the image registry, an **Failure to pull from quay** error message appears. The Red Hat OpenShift AI Operator might be unavailable for retrieval in the following circumstances:

- The image registry is unavailable.
- There is a problem with your network connection.
- Your cluster is not operational and is therefore unable to retrieve the image registry.

Diagnosis

Check the logs in the **Events** section in OpenShift for further information about the **Failure to pull from quay** error message.

Resolution

- Contact Red Hat support.

10.2. OPENSIFT AI DOES NOT INSTALL ON UNSUPPORTED INFRASTRUCTURE

Problem

You are deploying on an environment that is not documented as supported by the Red Hat OpenShift AI Operator.

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads** → **Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.

4. Click the **rhods-operator-`<random string>`** pod that shows an error in the **Status** column. The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for the **ERROR: Deploying on \$infrastructure, which is not supported. Failing Installation** error message.

Resolution

- Before proceeding with a new installation, ensure that you have a fully supported environment on which to install OpenShift AI. For more information, see [Supported Configurations for 3.x](#).

10.3. THE CREATION OF THE OPENSIFT AI CUSTOM RESOURCE (CR) FAILS

Problem

During the installation process, the OpenShift AI Custom Resource (CR) does not get created. This issue occurs in unknown circumstances.

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads** → **Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.
4. Click the **rhods-operator-`<random string>`** pod that shows an error in the **Status** column. The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for the **ERROR: Attempt to create the ODH CR failed.** error message.

Resolution

- Contact Red Hat support.

10.4. THE CREATION OF THE OPENSIFT AI NOTEBOOKS CUSTOM RESOURCE (CR) FAILS

Problem

During the installation process, the OpenShift AI Notebooks Custom Resource (CR) does not get created. This issue occurs in unknown circumstances.

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads → Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.
4. Click the **rhods-operator-*<random string>*** pod that shows an error in the **Status** column.
The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for the **ERROR: Attempt to create the RHODS Notebooks CR failed.** error message.

Resolution

- Contact Red Hat support.

10.5. THE OPENSIFT AI DASHBOARD IS NOT ACCESSIBLE

Problem

After installing OpenShift AI, the **redhat-ods-applications**, **redhat-ods-monitoring**, and **redhat-ods-operator** project namespaces are **Active** but you cannot access the dashboard due to an error in one of the pods.

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads → Pods**.
3. Set the **Project** to **All Projects**.
4. Click **Filter** and select the checkbox for every status except **Running** and **Completed**.
The page displays the pods that have an error.

Resolution

- To see more information and troubleshooting steps for a pod, on the **Pods** page, click the link in the **Status** column for the pod.
- If the **Status** column does not display a link, click the pod name to open the pod details page and then click the **Logs** tab.

10.6. REINSTALLING OPENSIFT AI FAILS WITH AN ERROR

Problem

After uninstalling the OpenShift AI Operator and reinstalling it by using the CLI, the reinstallation fails with an **unable to find DSCInitialization** error in one of the OpenShift AI Operator pod logs. This issue can occur if the **Auth** custom resource from the previous installation was not deleted after uninstalling

the OpenShift AI Operator and before reinstalling it. For more information, see [Understanding the uninstallation process](#).

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads** → **Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.
4. Click the **rhods-operator-<random string>** pod that shows an error in the **Status** column. The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for an error message similar to the following:

```
{ "name": "auth", "namespace": "", "name": "auth", "reconcileID": "7bff53ae-1252-46fe-831a-fdc824078a1b", "error": "unable to find DSCInitialization", "stacktrace": "sigs.k8s.io/controller-runtime/pkg/internal/controller." }
```

Resolution

1. Uninstall the OpenShift AI Operator.
2. Delete the **Auth** custom resource:
 - a. In the OpenShift web console, switch to the **Administrator** perspective.
 - b. Click **API Explorer**.
 - c. From the **All groups** drop-down list, select or enter **services.platform.opendatahub.io**.
 - d. Click the **Auth** kind.
 - e. Click the **Instances** tab.
 - f. Click the action menu (⋮) and select **Delete Auth**. The **Delete Auth** dialog appears.
 - g. Click **Delete**.
3. Install the OpenShift AI Operator again.

10.7. THE DEDICATED-ADMINS ROLE-BASED ACCESS CONTROL (RBAC) POLICY CANNOT BE CREATED

Problem

The Role-based access control (RBAC) policy for the dedicated-admins group in the target project cannot be created. This issue occurs in unknown circumstances.

Diagnosis

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads → Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.
4. Click the **rhods-operator-`<random string>`** pod that shows an error in the **Status** column.
The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for the **ERROR: Attempt to create the RBAC policy for dedicated admins group in \$target_project failed.** error message.

Resolution

- Contact Red Hat support.

10.8. THE ODH PARAMETER SECRET DOES NOT GET CREATED

Problem

An issue with the OpenShift AI Operator's flow could result in failure to create the ODH parameter.

Diagnosis

1. In the OpenShift web console, switch to the **Administrator** perspective.
2. Click **Workloads → Pods**.
3. Set the **Project** to **All Projects** or **redhat-ods-operator**.
4. Click the **rhods-operator-`<random string>`** pod that shows an error in the **Status** column.
The **Pod details** page appears.
5. Click **Logs**.
6. Select **rhods-operator** from the drop-down list.
7. Check the log for the **ERROR: Addon managed odh parameter secret does not exist.** error message.

Resolution

- Contact Red Hat support.

CHAPTER 11. UNINSTALLING RED HAT OPENSIFT AI SELF-MANAGED

This section shows how to use the OpenShift CLI (**oc**) to uninstall the Red Hat OpenShift AI Operator and any OpenShift AI components installed and managed by the Operator.



NOTE

Using the OpenShift CLI (**oc**) is the recommended way to uninstall the Operator. Depending on your version of OpenShift, using the web console to perform the uninstallation might not prompt you to uninstall all associated components. This could leave you unclear about the final state of your cluster.

11.1. UNDERSTANDING THE UNINSTALLATION PROCESS

Installing Red Hat OpenShift AI created several custom resource instances on your OpenShift cluster for various components of OpenShift AI. After installation, users likely created several additional resources while using OpenShift AI. Uninstalling OpenShift AI removes the resources that were created by the Operator, but retains the resources created by users to prevent inadvertently deleting information you might want.

What is deleted

Uninstalling OpenShift AI removes the following resources from your OpenShift cluster:

- **DataScienceCluster** custom resource instance and the custom resource instances it created for each component
- **DSCInitialization** custom resource instance
- **Auth** custom resource instance created during or after installation
- **FeatureTracker** custom resource instances created during or after installation
- **ServiceMesh** custom resource instance created by the Operator during or after installation
- **KNativeServing** custom resource instance created by the Operator during or after installation
- **redhat-ods-applications**, **redhat-ods-monitoring**, and **rhods-notebooks** namespaces created by the Operator
- Workloads in the **rhods-notebooks** namespace
- **Subscription**, **ClusterServiceVersion**, and **InstallPlan** objects
- **KfDef** object (version 1 Operator only)

What might remain

Uninstalling OpenShift AI retains the following resources in your OpenShift cluster:

- Projects created by users
- Custom resource instances created by users

- Custom resource definitions (CRDs) created by users or by the Operator

While these resources might still remain in your OpenShift cluster, they are not functional. After uninstalling, Red Hat recommends that you review the projects and custom resources in your OpenShift cluster and delete anything no longer in use to prevent potential issues, such as pipelines that cannot run, notebooks that cannot be undeployed, or models that cannot be undeployed.

Additional resources

[Operator Lifecycle Manager \(OLM\) uninstall documentation](#)

11.2. UNINSTALLING OPENSIFT AI SELF-MANAGED BY USING THE CLI

The following procedure shows how to use the OpenShift CLI (**oc**) to uninstall the Red Hat OpenShift AI Operator and any OpenShift AI components installed and managed by the Operator.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have installed the OpenShift CLI (**oc**) as described in the appropriate documentation for your cluster:
 - [Installing the OpenShift CLI](#) for OpenShift Container Platform
 - [Installing the OpenShift CLI](#) for Red Hat OpenShift Service on AWS
- You have backed up the persistent disks or volumes used by your persistent volume claims (PVCs).

Procedure

1. Open a new terminal window.
2. Log in to your OpenShift cluster as a cluster administrator, as shown in the following example:

```
$ oc login <openshift_cluster_url> -u system:admin
```

3. Create a **ConfigMap** object for deletion of the Red Hat OpenShift AI Operator.

```
$ oc create configmap delete-self-managed-odh -n redhat-ods-operator
```

4. To delete the **rhods-operator**, set the **addon-managed-odh-delete** label to **true**.

```
$ oc label configmap/delete-self-managed-odh api.openshift.com/addon-managed-odh-delete=true -n redhat-ods-operator
```

5. When all objects associated with the Operator are removed, delete the **redhat-ods-operator** project.

- a. Set an environment variable for the **redhat-ods-applications** project.

```
$ PROJECT_NAME=redhat-ods-applications
```

- b. Wait until the **redhat-ods-applications** project has been deleted.

```
while oc get project $PROJECT_NAME &> /dev/null; do
  echo "The $PROJECT_NAME project still exists"
  sleep 1
done
echo "The $PROJECT_NAME project no longer exists"
```

When the **redhat-ods-applications** project has been deleted, you see the following output.

```
The redhat-ods-applications project no longer exists
```

- c. When the **redhat-ods-applications** project has been deleted, delete the **redhat-ods-operator** project.

```
$ oc delete namespace redhat-ods-operator
```

Verification

1. Confirm that the **rhods-operator** subscription no longer exists.

```
$ oc get subscriptions --all-namespaces | grep rhods-operator
```

2. Confirm that the following projects no longer exist.

- **redhat-ods-applications**
- **redhat-ods-monitoring**
- **redhat-ods-operator**
- **rhods-notebooks**

```
$ oc get namespaces | grep -e redhat-ods* -e rhods*
```

The **rhods-notebooks** project existed only if you installed the workbenches component of OpenShift AI. See [Installing and managing Red Hat OpenShift AI components](#).